Joint Embedding of Graphs

Shangsi Wang, Carey E. Priebe, Joshua T. Vogelstein, Avanti Athreya, Minh Tang, Vince Lyzinski, Youngser Park

Abstract—In this paper, we propose a method to jointly embed multiple undirected graphs. The method identifies a linear subspace spanned by rank symmetric matrices and projects adjacency matrices of graphs into the subspace. The embedding coefficients can be treated as features of graphs. We propose a random graph model which can be used to model multiple random graphs. It can be shown that under the model our method produces estimates of parameters with small errors. We demonstrate through simulation experiments that our joint embedding method leads to state of the art performance in inference on multiple graphs. We further conduct a real data experiment which shows that individual creativity index can be predicted by the neural connectivity of human brain.

Index Terms—random graph, feature extraction, inference on graphs

✦

1 INTRODUCTION

IN many problems arising in science, graphs appear naturally as data to model complex relationship between a set of objects. Graphs have been used in various application domains as diverse as social networks [1], internet maps [2], brain connectomics [3], political voting networks [4], and many others. The graphs are naturally high dimensional objects with complicated topological structure which makes graph clustering and classification a challenge to traditional machine learning algorithms. Therefore, feature extraction and dimension reduction technique is in desperate need to study graphs. In this paper, we propose an algorithm to jointly embed multiple graphs into low dimensional space. We demonstrate through theories and experiments that our embedding algorithm produces features which lead to state of the art performance for subsequent inference tasks on graphs.

There exist a few unsupervised approaches to extract features from graphs. First, classical PCA can be applied by treating each edge of graphs a raw feature vector [5]. This approach produces features which are linear combinations of edges, but it ignores the topological structure of graphs and features extracted are not interpretable. Second, features can be extract by computing summary topological and label statistics from graphs [6]. These statistics commonly include number of edges, number of triangles, average clustering coefficient, maximum effective eccentricity, etc. In general, it is hard to know what intrinsic statistics to compute a priori and computing some statistics can be quite computational expensive. Also, some filtering criteria based on frequency of subgraph are proposed. For example, fast frequent subgraph mining algorithm can identify all connected subgraphs that occur in a large fraction of graphs in a graph data set [7]. Spectral feature selection can also be applied to graphs. It treats each graph sample as a node and constructed an object graph. Features are computed through the spectral decomposition of object graph [8].

The joint embedding takes a matrix factorization approach to extract features from graphs. The algorithm manages to simultaneously identify a set of rank one matrices and project adjacency matrices into the linear subspace spanned by the set of rank one matrices. Joint embedding can be understood as a generalization of adjacency spectral embedding for multiple graphs [9]. We demonstrate through simulation experiments that our joint embedding algorithm extracts features which lead to good performance for a variety of inference tasks. In the next section, we review some random graph models and present a model for generating multiple random graphs. In Section 3, we define the joint embedding of graphs and present an algorithm to compute it. In Section 4, we perform some theoretical analysis of joint embedding. The theoretical results and a real data experiment are explored in Section 5. We conclude the paper with a brief discussion of possible future work.

2 SETTING

In this paper, we focus on embedding unweighted and undirected graphs for simplicity, although the joint embedding algorithm works on weighted and directed graphs as well. Let $\{G_i = (V_i, E_i)\}_{i=1}^m$ be m graphs each with n vertices. We require the vertices in these graphs to be matched, which means that all the graphs have a common vertex set V. The joint embedding algorithm tries to embed all G_is simultaneously into $\mathbb{R}^d$ and represent G_i by a vector $\lambda_i \in \mathbb{R}^d$. Before discussing the joint embedding algorithm, we first introduce three random graph models.

Definition (Stochastic Block Model (SBM)) Let π be a prior probability vector for block membership. Denote $\tau = (\tau_1, \tau_2, ..., \tau_n) \in [K]^n$ be a block membership vector, where τ_i are independent and identically distributed according to probability vector $\pi \in \mathbb{R}^K$. Denote $B \in [0,1]^{K \times K}$ to be the block connecting probability matrix. Suppose A is a random adjacency matrix given by,

$$P(A|\tau, B) = \prod_{i<j} B_{\tau_i,\tau_j}^{A_{i,j}}(1 - B_{\tau_i,\tau_j})^{(1-A_{i,j})}$$

Then, we say A is an adjacency matrix of a K-block stochastic block model graph, and we write $A \sim SBM(\pi, B)$. We may also treat τ as the parameter of interest, then we write $A \sim SBM(\tau, B)$.

SBM is a widely used model to study community structure of a graph [10]. We then recall the notion of a random dot product graph [11].

Definition (Random Dot Product Graph (RDPG)) Let F be a distribution on a set $\mathcal{X} \in \mathbb{R}^d$ satisfying $x^T y \in [0,1]$ for all $x, y \in \mathcal{X}$. Let $X = [X_1^T, X_2^T, ..., X_n^T] \in \mathbb{R}^{n \times d}$. We say $(X, A) \sim RDPG(F)$ if X_i is independent and identically distributed according to F and condition on X, A_{ij} is an independent Bernoulli random variable,

$$A_{ij} \sim Bernoulli(X_i^T X_j).$$

Alternatively,

$$P(A|X) = \prod_{i<j} X_i^T X_j^{A_{ij}} (1 - X_i^T X_j)^{1-A_{ij}}$$

Also, we define $P := XX^T$ to be edge probability matrix. When we are more interested in estimating latent positions X, we treat X as parameters and write $A \sim RDPG(X)$.

Definition (Multiple Random Eigen Graphs (MREG)) Let $\{h_k\}_{k=1}^d$ be a set of norm-1 vectors in $\mathbb{R}^n$, and F be a distribution on a set $\mathcal{X} \in \mathbb{R}^d$, satisfying $\sum_{k=1}^d \lambda[k] h_k h_k^T \in [0,1]^{n \times n}$ for all $\lambda \in \mathcal{X}$, where $\lambda[k]$ is the kth entry of vector λ. We say $\{(\lambda_i, A_i)\}_{i=1}^m$ follows a m-graph d-dimensional multiple random eigen graphs model and write

$$\{(\lambda_i, A_i)\}_{i=1}^m \sim MREG(F, h_1, ..., h_d)$$

if λ_i is independent and identically distributed according to F, and condition on λ_i, the entries of A_i follows independent Bernoulli distribution,

$$A_i[s,t] \sim Bernoulli(\sum_{k=1}^d \lambda_i[k] h_k[s] h_k[t]).$$

We call $P_i := \sum_{k=1}^d \lambda_i[k] h_k h_k^T$ the edge probability matrix for graph i. In cases that we are more interested in $\{\lambda_i\}_{i=1}^m$, we treat them as parameters and write

$$\{A_i\}_{i=1}^m \sim MREG(\lambda_1, ..., \lambda_m, h_1, ..., h_d).$$

In MREG, we allow self loops to happen. It is mainly for the convenience of proving theories. If we ignore the self loops, we can have the following relationships between three random graph models. If an adjacency matrix $A \sim SBM(\pi, B)$ and the block connecting matrix B is semi-positive definite, A can also be written as a $RDPG(F)$ with F being a finite mixture of point masses. If an adjacency matrix $A \sim RDPG(X)$, then it is also a 1-graph $MREG(\lambda_1, h_1, ..., h_d)$ with h_k being the normalized kth column of X and λ_1 being the vector containing squared norm of kth column of X. However, 1-graph $MREG(\lambda_1, h_1, ..., h_d)$ is not necessarily a RDPG graph since λ_1 could contain negative entries which may result in indefinite edge probability matrix. If m RDPG graphs have edge probability matrices sharing the same eigenspace, they also follow a m-graph MREG model.

3 METHODOLOGY

3.1 Joint Embedding of Graphs

The joint embedding method considers a collection of vertex-aligned graphs, and estimates a common embedding space across all graphs and a loading for each graph. Specifically, it simultaneously identifies a subspace spanned by a set of rank one symmetric matrices $\{\hat{h}_k \hat{h}_k^T\}_{k=1}^d$ and projects each adjacency matrix A_i linearly into the subspace. The coefficients obtained by projecting A_i is denoted by $\hat{\lambda}_i \in \mathbb{R}^d$ which is called loadings for graph i. To estimate rank one symmetric matrices and loadings for graphs, we minimize the sum of squared Frobenius distances between adjacency matrices and their projections. The detailed definition of joint embedding of graphs is provided below.

Definition (Joint Embeddings of Graphs) Given m graphs $\{G_i\}_{i=1}^m$ with n vertices, let $A_i \in \{0,1\}^{n \times n}$ denote the adjacency matrices of graph G_i. The d-dimensional joint embedding of graphs $\{G_i\}_{i=1}^m$ are given by

$$(\hat{\lambda}_1, ..., \hat{\lambda}_m, \hat{h}_1, ..., \hat{h}_d) = \underset{\lambda_i, \|h_k\|=1}{\operatorname{argmin}} \sum_{i=1}^m \|A_i - \sum_{k=1}^d \lambda_i[k] h_k h_k^T\|^2.$$
(1)

Here, $\|\cdot\|$ denotes the Frobenius norm and $\lambda_i[k]$ is the kth entry of vector λ_i

To make the model identifiable, we put an extra requirement on the norm of h_k. To ease our notation, we introduce two matrices $\Lambda \in \mathbb{R}^{m \times d}$ and $H \in \mathbb{R}^{n \times d}$, where λ_i is the ith row of Λ and h_k is the kth row of H. That is $\Lambda = [\lambda_1^T, ..., \lambda_m^T]$ and $H = [h_1, ..., h_d]$. We can rewrite equation (1) using Λ and H,

$$(\hat{\Lambda}, \hat{H}) = \underset{\Lambda, \|h_k\|=1}{\operatorname{argmin}} \sum_{i=1}^m \|A_i - \sum_{k=1}^d \Lambda_{ik} h_k h_k^T\|^2.$$

We denote the function on the left hand side of the equation by $f(\Lambda, H)$ which is implicitly a function of λ_is and h_ks. There are several alternative ways to formulate the problem. If we convert λ_i into a diagonal of matrix $D_i \in \mathbb{R}^{d \times d}$ by putting entries of λ_i on the diagonal of D_i, then solving equation (1) is equivalent to

$$\underset{D_i, \|h_k\|=1}{\operatorname{argmin}} \sum_{i=1}^m \|A_i - HD_iH^T\|^2, \text{ subject to } D_i \text{ is diagonal}.$$

We may also view (1) as a tensor factorization problem. If $\{A_i\}_{i=1}^m$ are stacked up to a 3-D array $\mathbf{A} \in \mathbb{R}^{m \times n \times n}$, then solving equation (1) is also equivalent to

$$\underset{\Lambda, \|h_k\|=1}{\operatorname{argmin}} \|\mathbf{A} - \sum_{k=1}^d \Lambda_{*k} \otimes h_k \otimes h_k\|^2$$

where $\otimes$ denotes the tensor product and Λ_{*k} is the kth column of Λ.

The optimization problem in equation (1) is also similar to principle component analysis in the sense of minimizing squared reconstruction error to recover loadings and components. However, there are extra symmetries and rank constrains on the components. In case of embedding only one graph, the joint embedding is equivalent to the adjacency spectral embedding solved by singular value decomposition. Next, we describe an algorithm to optimize the objective function $f(\Lambda, H)$.

3.2 Alternating Descent Algorithm

Joint embedding of $\{G_i\}_{i=1}^m$ can be estimated by solving the optimization problem in equation (1). There are a few methods proposed to solve similar problems. Carroll and Chang [12] propose to use an alternating method that ignores symmetry, with the idea that it will often converge to a symmetric solution. Gradient approaches have also been considered for similar problems [13] [14]. We develop an alternating descent algorithm to minimize $f(\Lambda, H)$. The algorithm iteratively updates one of Λ and H while treating the other parameter as fixed. We will see that optimizing Λ when fixing H is easy, since it is essentially a least square problem. However, optimizing H when fixing Λ is hard due to the fact that the problem is non-convex and there is no closed form solution available. In this case, we use gradient information and take an Armijo line search strategy to update H [15].

Instead of optimizing all columns Λ and H simultaneously, we consider a greedy algorithm which solves the optimization problem by only considering one column of Λ and H at a time. Specifically at iteration k_0, we fix all estimates for first $k_0 - 1$ columns of Λ and H, and then the objective function is minimized by searching through kth column of Λ and H, that is

$$(\hat{\Lambda}_{*k_0}, \hat{h}_{k_0}) = \operatorname*{argmin}_{\Lambda_{*k_0}, \|h_{k_0}\|=1}$$
$$\sum_{i=1}^m \|A_i - \sum_{k=1}^{k_0-1} \hat{\Lambda}_{ik}\hat{h}_k\hat{h}_k^T - \Lambda_{ik_0}h_{k_0}h_{k_0}^T\|^2 \quad (2)$$

We denote the function on the left hand side of the equation by $f(\Lambda_{*k_0}, h_{k_0})$. To compute a d-dimensional joint embedding $(\hat{\Lambda}, \hat{H})$, we recursively solve the one dimensional optimization problem above by letting k_0 to vary from 1 to d.

There are a few advantages in recursively solving one dimensional problem. First, there are less parameters to fit at each iteration, since we are only allowed to vary Λ_{*k_0} and h_{k_0}. It makes initialization and optimization procedure much easier compared to optimizing all columns of H at the same time. Second, it implicitly enforces an ordering on columns of H. The ordering allows us to select the top few columns of Λ and H in cases we want to perform model selection. Third, it allows incremental computation. If we fit d and d' dimensional joint embeddings, the first $\min(d, d')$ columns of $\hat{\Lambda}$ and $\hat{H}$ will be the same. The disadvantage of optimizing each column separately is that we are more likely to end up at a local minimum when the objective function is structured not in favor of embedding separately. In practice, this problem can be mitigated by running the joint embedding algorithm several times with random initialization.

To update Λ_{*k_0} and h_{k_0} in one dimensional problem, we need to evaluate two derivatives: $\frac{\partial f}{\partial h_{k_0}}$ and $\frac{\partial f}{\partial \Lambda_{ik_0}}$. Denote the residual matrix $R_{ik_0} = A_i - \sum_{k=1}^{k_0-1} \hat{\Lambda}_{ik}\hat{h}_k\hat{h}_k^T$. The gradient

of objective with respect to h_{k_0} is given by,

$$\frac{\partial f}{\partial h_{k_0}} = -4\sum_{i=1}^m \Lambda_{ik_0}(R_{ik} - \Lambda_{ik_0}h_{k_0}h_{k_0}^T)h_{k_0} \quad (3)$$

The derivative of objective with respect to Λ_{ik_0} are given by,

$$\frac{\partial f}{\partial \Lambda_{ik_0}} = -2\langle R_{ik} - \Lambda_{ik_0}h_{k_0}h_{k_0}^T, h_{k_0}h_{k_0}^T\rangle$$

Set the derivative to 0, we have

$$\hat{\Lambda}_{ik_0} = \langle R_{ik}, h_{k_0}h_{k_0}^T\rangle \quad (4)$$

Algorithm 1 describes the general procedure to compute d-dimensional joint embedding of graphs $\{G_i\}_{i=1}^m$.

Algorithm 1 Joint Embedding Algorithm

1: **procedure** FIND JOINT EMBEDDINGS $\hat{\Lambda}, \hat{H}$ OF $\{A_i\}_{i=1}^m$
2:     Set residuals: $R_{i1} = A_i$
3:     **for** $k = 1 : d$ **do**
4:         Initialize h_k and Λ_{*k}
5:         **while** not convergent **do**
6:             Fixing Λ_{*k}, update h_k by gradient descent (3)
7:             Fixing h_k, update Λ_{*k} by (4)
8:             Compute objective $\sum_{i=1}^m \|R_{ik} - \Lambda_{ik}h_kh_k^T\|^2$
9:         **end while**
10:        Update residuals: $R_{i(k+1)} = R_{ik} - \Lambda_{ik}h_kh_k^T$
11:     **end for**
12:     Output $\hat{\Lambda} = [\Lambda_{*1}, ..., \Lambda_{*d}]$ and $\hat{H} = [h_1, ..., h_d]$
13: **end procedure**

After the algorithm finish, rows of $\hat{\Lambda}$ denoted by $\{\hat{\lambda}_i\}_{i=1}^m$ are treated as features for graphs. If a new graph G is observed with adjacency matrix A, we can project A linearly onto matrices $\{\hat{h}_k\hat{h}_k^T\}_{k=1}^d$ to obtain features for the graph. We should admit that the optimization algorithm may not be the best approach to solve the problem; however, it is not the focus of this paper. Based on results from numerical applications, it works well in estimating parameters and extracting features for subsequent inference.

3.3 Variations

The joint embedding algorithm described above can be modified to accommodate several settings.

Variation 1. When all graphs come from the same model, we can force Λ_{ik}s to be equal across graphs. This is useful when primary inference task is to extract features for vertices. Since all graphs share the same loadings, with slightly abusing notations, let Λ be a vector in $\mathbb{R}^d$ and the optimization problem becomes,

$$(\hat{\Lambda}, \hat{H}) = \operatorname*{argmin}_{\Lambda, \|h_k\|=1} \sum_{i=1}^m \|A_i - \sum_{k=1}^d \Lambda_k h_k h_k^T\|^2$$

In this case, the optimization problem can be solved exactly by finding singular value decomposition of the average adjacency matrix $\frac{1}{m}\sum_{i=1}^m A_i$.

Variation 2. When there is a discrete label $y_i \in \mathbb{Y}$ associated with G_i available, we may require all Λ_{ik}s to be equal within

class, that is $\Lambda_{ik} = \Lambda_{y_i k}$ for all i. If we let $\Lambda \in \mathbb{R}^{|\mathbb{Y}| \times d}$, the optimization problem becomes,

$$(\hat{\Lambda}, \hat{H}) = \underset{\Lambda, \|h_k\|=1}{\operatorname{argmin}} \sum_{i=1}^{m} \|A_i - \sum_{k=1}^{d} \Lambda_{y_i k} h_k h_k^T\|^2$$

In this case, when updating Λ in algorithm 1, we should average $\Lambda_{y_i k}$ within the same class.

Variation 3. We may require all Λ_{ik}s to be greater than 0 like in the non-negative matrix factorization. The advantage of this constrain is that graph G_i may be automatically clustered based on largest entry of $\hat{\lambda}_i$.

$$(\hat{\Lambda}, \hat{H}) = \underset{\Lambda \geq 0, \|H_{*k}\|=1}{\operatorname{argmin}} \sum_{i=1}^{m} \|A_i - \sum_{k=1}^{d} \Lambda_{ik} H_{*k} H_{*k}^T\|^2$$

Next, we discuss properties of joint embedding when treated as a parameter estimation procedure for MREG model.

4 THEORY

In this section, we consider a simple setting where graphs follow 1-dimensional MREG model, that is $\{(\lambda_i, A_i)\}_{i=1}^{m} \sim MREG(F, h_1)$. Under this model, we can understand joint embedding of graphs as an estimator. $\hat{\lambda}_i$ and $\hat{h}_1$ are estimates of λ_i and h. We prove the following two theorems concerning the asymptotic behavior of estimator $\hat{h}_1$ produced by joint embedding. Let $\hat{h}_1^m$ denote the estimates based on m graphs and define functions ρ, D_m and D by

$$\rho(A_i, h) = \|A_i - \langle A_i, hh^T \rangle hh^T\|^2$$

$$D_m(h, h_1) = \frac{1}{m} \sum_{i=1}^{m} \rho(A_i, h)$$

$$D(h, h_1) = E(\rho(A_i, h))$$

By equation (1), we have

$$\hat{h}_1^m = \underset{\|h\|=1}{\operatorname{argmin}} \underset{\lambda_i}{\operatorname{argmin}} \sum_{i=1}^{m} \|A_i - \lambda_i hh^T\|$$

By equation (4), we have

$$\langle A_i, hh^T \rangle = \underset{\lambda_i}{\operatorname{argmin}} \sum_{i=1}^{m} \|A_i - \lambda_i hh^T\|$$

Therefore,

$$\hat{h}_1^m = \underset{\|h\|=1}{\operatorname{argmin}} D_m(h, h_1)$$

The first theorem states that $\hat{h}_1^m$ converges almost surely to the global minimum of $D(h, h_1)$ given the global minimum is unique.

Theorem 4.1. If $D(h, h_1)$ has a unique global minimum at h', then $\hat{h}_1^m$ converges almost surely to h' as m goes to infinity, that is

$$\hat{h}_1^m \overset{a.s.}{\to} h'$$

In theorem 1, we require h' to be the unique global minimizer of $D(h, h_1)$. However, the global optimizer is definitely not unique due to the symmetry up to sign flip of h, that is $D(h, h_1) = D(-h, h_1)$ for any h. This problem can be fixed by forcing an orientation of $\hat{h}_1^m$ or stating that the

convergence is up to a sign flip. It is also possible that there are multiple global minimizers of $D(h, h_1)$ which are not sign flip of each other. In this case, theorem 1 does not apply. We are currently only certain that when all graphs are from Erdos-Renyi random graph model, the global minimizer is unique up to a sign flip. The next theorem concerns the asymptotic bias of $\hat{h}_1^m$.

Theorem 4.2. If h' is a minimizer of $D(h, h_1)$, then

$$\|h' - h_1\| \leq \frac{2E(\lambda_i)}{E(\lambda_i^2)(h_1^T h')^2}$$

To see an application of theorem 4.2, we consider cases that all graphs are Erdos-Renyi graphs with 100 vertices and edge probability 0.5. Under this setting, theorem 2 implies $\|h' - h_1\| \in [0, 0.04] \cup [1.28, 1.52]$. The second interval is disturbing. It is due to the fact that when $h_1^T h'$ is small the bound is useless. We provide some insights why the second interval is there and how we can get rid of it with more assumptions. In the proof, we show that

$$h' = \underset{\|h\|=1}{\operatorname{argmax}} E(\langle A_i, hh^T \rangle^2)$$

If we take a closer look on $E(\langle A_i, hh^T \rangle^2)$,

$$E(\langle A_i, hh^T \rangle^2) = E(\langle P_i, hh^T \rangle^2) + E(\langle A_i - P_i, hh^T \rangle^2)$$
$$= E(\lambda_i^2)(h_1^T h)^4 + E((h^T(A_i - P_i)h)^2)$$

Therefore,

$$h' = \underset{\|h\|=1}{\operatorname{argmax}} E(\lambda_i^2)(h_1^T h)^4 + E((h^T(A_i - P_i)h)^2)$$

We see $E(\lambda_i^2)(h_1^T h)^4$ is clearly maximized when $h = h_1$; however, the noise term $E((h^T(A_i - P_i)h)^2)$ is generally not maximized at $h = h_1$. If we assume n is large, we can apply concentration inequality to $(h^T(A_i - P_i)h)^2$ and have an upper bound on $E((h^T(A_i - P_i)h)^2)$. If we further assume A_i is not too sparse, that is $E(\lambda_i^2)$ grows with n fast enough, then the sum of these two terms are dominated by the first term. This provides a way to have a lower bound on $h_1^T h'$. We may then remove the second interval. In general, if n is small, the noise term may cause h' differs with h_1 by a significant amount. We focus on the case that n is fixed in this paper. The case that n goes to infinity for random inner product graph is considered in [16].

The two theorems above only concern the estimation of h_1, but not λ_i. Based on equation (4), we estimate λ_i by

$$\hat{\lambda}_i^m = \langle A_i, \hat{h}_1^m \hat{h}_1^{mT} \rangle$$

If n is fixed, it is impossible to prove any consistency result on estimation of λ_i due to the fact we only observe one finite graph A_i associated with λ_i. When m goes to infinity, we can apply theorem 1,

$$\hat{\lambda}_i^m \to \langle A_i, h'h'^T \rangle = h'^T A_i h'$$

Then applying theorem 2 and utilizing the fact that $h^T A_i h$ is continuous in h, we may have a upper bound on $|\hat{\lambda}_i^m - h_1^T A_i h_1|$. In general, $h_1^T A_i h_1$ is concentrated about λ_i with high probability.

5 EXPERIMENTS

Before going into details of experiments, we want to discuss how to select the dimensionality of embedding d. Estimating d is an important model selection question which has been studied for years under various settings. It is not the focus of this paper, but we still face this decision in numerical experiments. In the simulation experiments of this section, we assume d is known to us and simply set our dimension estimate $\hat{d}$ equal to the d. In the real data experiment, we recommend two approaches to determine $\hat{d}$. Both approaches require first run d'-dimensional joint embedding algorithm, where d' needs to be large and we are confident d is less d'. We then plot the objective function versus dimension, and determine $\hat{d}$ where the objective start to flatten out. Alternatively, we can plot $\{\hat{\Lambda}_{ik}\}_{i=1}^m$ for $k = 1, ...d'$, and select $\hat{d}$ where the average loadings start to look like noise with 0 mean. These two approaches usually yield similar dimensionality estimates $\hat{d}$.

5.1 Simulation Experiment 1: Joint Embedding in Simple Settings

Here we present a numerical example demonstrating some properties of our joint embeddings as the number of graphs grows. We repeatedly generate graphs with 5 vertices from 1-dimensional MREG where $\lambda_i \sim Unif(1,2)$ and $h_1 = [0.1, 0.2, 0.3, 0.4, 0.5]^T/0.74$. We keep doubling the number of graphs m as it increases from 2^4 to 2^{16}. At each value of m, we compute 1-dimensional joint embedding of m graphs. Let the estimated parameters based on m graphs denote by $\hat{\lambda}_i^m$ and $\hat{h}_1^m$. Two quantities based on $\hat{h}_1^m$ are calculated. The first one is the norm difference between the current h_1 estimates and the previous estimates, namely $\|\hat{h}_1^m - \hat{h}_1^{m/2}\|$. It tries to seek for numerical evidences for the convergence of our principled estimation procedure. The second quantity is $\|\hat{h}_1^m - h_1\|$. It tries to see whether $\hat{h}_1$ is an unbiased estimator for h_1. The plot at bottom demonstrates the result.


Fig. 1. Bias and Convergence of Joint Embedding

From the plot, we can see that the norm difference converges to 0 as m increases. It suggests the convergence of $\hat{h}_1^m$. Secondly, we should notice that the bias $\|\hat{h}_1^m - h_1\|$ does not converge to 0; instead, it stops to decrease at around 0.03. This suggests that $\hat{h}_1$ is an asymptotically biased estimator for h_1. The $\hat{h}_1^m$ with $m = 2^{16}$ is

$[0.083, 0.186, 0.296, 0.406, 0.503]^T/0.74$. Compared to h_1, there are negative biases for small entries of h_1, and positive biases for large entries of h_1. Actually, there are some theoretical evidences for this phenomenon as well. In fact, when there are infinitely many nuisance parameters present, Neyman and Scott gives an example that maximum likelihood estimator can be inconsistent as well [17]. In our case, there are infinitely many λ_is; therefore, we do not expect joint embedding to be an consistent estimator.

In applications like clustering or classifying multiple graphs, we may be not interested in $\hat{h}_1$. We are more interested in $\hat{\lambda}_i$ which provides information specifically about graph G_i. Here we consider two approaches to estimate λ_i. The first approach is estimating λ_i through joint embedding,

$$\hat{\lambda}_i = \langle A_i, \hat{h}_1^m \hat{h}_1^{mT} \rangle$$

The second approach estimates λ_i by forcing $\hat{h}_1$ equal to h_1, that is

$$\hat{\lambda}_i = \langle A_i, h_1 h_1^T \rangle$$

$\hat{\lambda}_i$ calculated this way can be thought of as 'oracle' estimates, since it assumes h_1 is known. For each value of m, we provide a boxplot shows the differences in estimates provided by two approaches. Not surprisingly, the differences are small due to the fact that $\hat{h}_1^m$ and h_1 are close.


Fig. 2. Boxplot for Lambda Estimates

5.2 Simulation Experiment 2: Joint Embedding to Classify Graphs

In this experiment, we consider the inference task of classifying graphs. The set up is m pairs $\{(A_i, Y_i)\}_{i=1}^m$ are observed. Each pair consists of an adjacency matrix $A_i \in \{0,1\}^{n \times n}$ and a label $Y_i \in [K]$. Furthermore, all pairs are assumed to be independent and identically distributed according to an unknown distribution $\mathbb{F}_{A,Y}$, that is

$$(A_1, Y_1), (A_2, Y_2), ..., (A_m, Y_m) \overset{i.i.d.}{\sim} \mathbb{F}_{A,Y}$$

The goal is to find a classifier g which is a function $g : \{0,1\}^{n \times n} \to [K]$ which has small classification error $L_g = P(g(A) \neq Y)$.

We consider a binary classification problem where Y can only take value 0 or 1. We independently generate 200

graphs with 100 vertices from two MREG models each. Let h_1 and h_2 be two vectors in $\mathbb{R}^{100}$, and

$$h_1 = [0.1, ..., 0.1]^T \text{, and } h_2 = [-0.1, ..., -0.1, 0.1, ..., 0.1]^T$$

Here, h_2 has -0.1 as its first 50 entries and 0.1 as its last 50 entries. Conditioned on the value of Y, we generate graphs according to two MREG models, that is

$$A_i | Y_i = 0 \sim MREG(F_0, h_1, h_2) \text{, with } F_0 = \mathbb{1}_{[25,5]}$$

$$A_i | Y_i = 1 \sim MREG(F_1, h_1, h_2) \text{, with } F_1 = \mathbb{1}_{[22.5,2.5]}$$

In terms of SBM, this graph generation scheme is also equivalent to,

$$A_i | Y_i = 0 \sim SBM\left((1, ..., 1, 2, ..., 2), \begin{bmatrix} 0.3 & 0.2 \\ 0.2 & 0.3 \end{bmatrix}\right)$$

$$A_i | Y_i = 1 \sim SBM\left((1, ..., 1, 2, ..., 2), \begin{bmatrix} 0.25 & 0.2 \\ 0.2 & 0.25 \end{bmatrix}\right)$$

To classify graphs, we first jointly embed 200 graphs. The embeddings are shown in the figure below. We can see two classes are clearly separated after being jointly embedded. Then, an 1-nearest neighbor classifier is constructed based


Fig. 3. Λ_i of jointly embedded graphs

on loadings $\{\hat{\lambda}_i\}_{i=1}^m$. The 1-NN classification error with 200 graphs is 0.

We compare classification performance using joint embedding and laplacian eigenmap [18]. For laplacian eigenmap, we first embed each normalized laplacian matrix and then compute Frobenious norm between embeddings. We also apply a 1-NN rule to classify graphs. We let the number of graphs m to increase from 4 to 200. For each value of m, we repeat the simulation 100 times. The result is shown in the next figure. We see joint embedding takes advantage of increasing sample size and dominates ASE at all values of m.

5.3 Real Data Experiment: CCI

In this experiment, we consider predicting individual composite creativity index (CCI) through functional (structural???) magnetic resonance imaging. Neuroimaging and creativity have been jointly investigated previously [19]. Most of studies utilize a statistical testing method and find CCI significant related or inversely related to the activity of some regions of brain. We embrace a different approach by


Fig. 4. Classification Performance of Joint Embedding and ASE

finding a regression model for CCI. We first jointly embed brain images and then regress CCI loadings estimated. 113 healthy, young adult subjects were scanned on using (???) scanner with functional magnetic resonance imaging. The image is then registered by Desikan atlas (???) [20] and a graph of 70 vertices is constructed through computing pairwise correlations across regions. DT measures were scored by (???) independent judges using the Consensual Assessment Technique [21], from which a CCI was derived. To predict CCI, we jointly embed 113 graphs with $d = 10$. Then, we construct a linear regression model by treating CCI as the response variable and $\hat{\lambda}_i$ as predictors, that is

$$CCI_i \sim \beta_0 + \lambda_i^T \beta + \epsilon_i$$

If only $\hat{\lambda}_i[1]$ is used as predictor, the figure 5 shows the result. We can see a linear relationship between CCI and first dimension loadings. If fit CCI to all the loadings, a summary of the linear model is provided in appendix. The R-square is 0.2325 and the model is statistically significant better compared to the null model with a p-value 0.0018.


Fig. 5. Predicting Creativity Index with Joint Embeddings

6 CONCLUSION

In summary, we develop a joint embedding method which can simultaneously embed multiple graphs into low dimensional space. Joint embedding can be used to estimate simple and intuitive statistics for inference problems on multiple vertex matched graphs. Learning on multiple graphs has significant applications in diverse fields and our results have both theoretical and practical implications for the problem. As the real data experiment illustrate, joint

embedding is a practically viable inference procedure.

We also propose a Multiple Random Eigen Graphs model. It can be seen as a generalization of random dot product graph model or stochastic block model for multiple random graphs. We analyze the performance of joint embedding on this model under simple settings. It can be shown our method provides estimates with bounded error.

Our approach is intimately related other matrix factorization approaches like singular value decomposition and tensor factorization. Indeed, they all try to estimate a low dimensional representation of high dimensional objects. We are currently investigate joint embedding with more or less regularizations on parameters under difference real set ups.

APPENDIX

Proof of Theorem 4.1 First, we are going to show that $|D_n(h, h_1) - D(h, h_1)|$ converges uniformly to 0. We notice three facts,

(1) The set $\{h : \|h\| = 1\}$ is a compact set.
(2) The function $\rho(A_i, h)$ is continuous for all h.
(3) For all h, $\rho(A_i, h)$ is bounded by n^2.

Therefore, by the uniform law of large numbers [22], we have

$$\sup_h |D_m(h, h_1) - D(h, h_1)| \overset{a.s.}{\to} 0$$

To prove the claim of theorem, we use a similar technique employed by Bickel and Doksum [23]. By definition, we have $D_m(\hat{h}_1^m, h) \leq D_m(h', h)$ and $D(h', h) \leq D(\hat{h}_1^m, h)$. Using these two inequalities we have

$$D_m(h', h) - D(h', h) \geq D_m(\hat{h}_1^m, h) - D(h', h)$$
$$\geq D_m(\hat{h}_1^m, h) - D(\hat{h}_1^m, h)$$

Therefore,

$$|D_m(\hat{h}_1^m, h) - D(h', h)| \leq \max(|D_m(h', h) - D(h', h)|,$$
$$|D_m(\hat{h}_1^m, h) - D(\hat{h}_1^m, h)|)$$

This implies

$$|D_m(\hat{h}_1^m, h) - D(h', h)| \leq \sup_h |D_m(h, h_1) - D(h, h_1)|$$

Hence, $|D_m(\hat{h}_1^m, h) - D(h', h)|$ must converge almost surely to 0 ,that is

$$|D_m(\hat{h}_1^m, h) - D(h', h)| \overset{a.s.}{\to} 0$$

If $\hat{h}_1^m$ does not converge almost surely h', $\|\hat{h}_1^m - h'\| \geq \epsilon$ for some ϵ and infinitely many values of m. Since h' is the unique global minimum, $|D(\hat{h}_1^m, h) - D(h', h)| > \epsilon'$ for infinitely many values of m. This contradicts with the previous equation. Therefore, $\hat{h}_1^m$ must converges almost surely h'.

Proof of Theorem 4.2 The first part of proof is showing that h' is the eigenvector corresponds to the largest eigenvalue of $E(\langle A_i, h'h'^T \rangle A_i)$. Then, we show $E(\langle A_i, h'h'^T \rangle A_i)$ is close

to a properly scaled $E(P_i)$. To conclude, we apply Davis-Kahan theorem to the top eigenvector of both matrices and get the desired result. First, we notice that

$$\min_{\|h\|=1} D(h, h_1) = \min_{\|h\|=1} E(\|A_i - \langle A_i, hh^T \rangle hh^T\|^2)$$
$$= \min_{\|h\|=1} E(\langle A_i, A_i \rangle - \langle A_i, hh^T \rangle^2)$$
$$= E(\langle A_i, A_i \rangle) - \max_{\|h\|=1} E(\langle A_i, hh^T \rangle^2)$$

Therefore,

$$h' = \underset{\|h\|=1}{\operatorname{argmin}} D(h, h_1) = \underset{\|h\|=1}{\operatorname{argmax}} E(\langle A_i, hh^T \rangle^2) \quad (5)$$

Taking the derivative of $E(\langle A_i, hh^T \rangle^2) + c(h^T h - 1)$ with respect to h, we have

$$\frac{\partial E(\langle A_i, hh^T \rangle^2) + c(h^T h - 1)}{\partial h} = E(\frac{\partial \langle A_i, hh^T \rangle^2}{\partial h}) + 2ch$$
$$= 4E(\langle A_i, hh^T \rangle A_i)h + 2ch$$

Set it to 0,

$$E(\langle A_i, h'h'^T \rangle A_i)h' = -\frac{1}{2}ch'$$

Using the fact that, $\|h'\| = 1$, we can solve c

$$h'^T E(\langle A_i, h'h'^T \rangle A_i)h' = E(\langle A_i, h'h'^T \rangle^2) = -\frac{1}{2}c$$

Then, substitute c

$$E(\langle A_i, h'h'^T \rangle A_i)h' = E(\langle A_i, h'h'^T \rangle^2)h' \quad (6)$$

Therefore, we see h' is an eigenvector of $E(\langle A_i, h'h'^T \rangle A_i)$ and the corresponding eigenvalue is $E(\langle A_i, h'h'^T \rangle^2)$. Furthermore, $E(\langle A_i, h'h'^T \rangle^2)$ must be the eigenvalue with the largest magnitude. Assume not, then there exists a h'' with norm 1 such that

$$|h''^T E(\langle A_i, h'h'^T \rangle A_i)h''| = |E(\langle A_i, h'h'^T \rangle \langle A_i, h''h''^T \rangle)|$$
$$> E(\langle A_i, h'h'^T \rangle^2)$$

However, by Cauchy-Schwarz inequality we must have

$$E(\langle A_i, h''h''^T \rangle^2)E(\langle A_i, h'h'^T \rangle^2) >$$
$$|E(\langle A_i, h'h'^T \rangle \langle A_i, h''h''^T \rangle)|^2$$

This implies $E(\langle A_i, h''h''^T \rangle^2) > E(\langle A_i, h'h'^T \rangle^2)$ which contradicts equation (5). This concludes that h' is the eigenvector corresponds to the largest eigenvalue of $E(\langle A_i, h'h'^T \rangle A_i)$.

Next, we compute $E(\langle A_i, h'h'^T \rangle A_i)$.

$$E(\langle A_i, h'h'^T \rangle A_i | P_i)$$
$$= E(\langle A_i - P_i, h'h'^T \rangle (A_i - P_i) | P_i)$$
$$\quad + E(\langle A_i - P_i, h'h'^T \rangle P_i) | P_i)$$
$$\quad + E(\langle P_i, h'h'^T \rangle (A_i - P_i) | P_i) + E(\langle P_i, h'h'^T \rangle P_i | P_i)$$
$$= E(\langle A_i - P_i, h'h'^T \rangle (A_i - P_i) | P_i) + \lambda_i (h_1^T h')^2 P_i$$
$$= 2h'h'^T * P_i * (J - P_i) - DIAG(h_1 h_1^T * P_i * (J - P_i))$$
$$\quad + \lambda_i (h_1^T h')^2 P_i$$

Here, DIAG() means only keep the diagonal of matrix; $*$ means the Hadamard product, and J is a matrix of all ones. Using the fact that $P_i = \lambda_i h_1 h_1^T$, we have

$$E(\langle A_i, h'h'^T \rangle A_i) - E(\lambda_i^2)(h_1^T h')^2 h_1 h_1^T$$
$$= E(E(\langle A_i, h'h'^T \rangle A_i | P_i) - \lambda_i (h_1^T h')^2 P_i)$$
$$= E(2h'h'^T * P_i * (J - P_i) - DIAG(h'h'^T * P_i * (J - P_i)))$$

If we consider the norm difference between $E(\langle A_i, h'h'^T \rangle A_i)$ and $E(\lambda_i^2)(h_1^T h')^2 h_1 h_1^T$, we have

$$\|E(\langle A_i, h'h'^T \rangle A_i) - E(\lambda_i^2)(h_1^T h')^2 h_1 h_1^T\|$$
$$= \|E(2h'h'^T * P_i * (J - P_i) - DIAG(h'h'^T * P_i * (J - P_i)))\|$$
$$\leq E(\|2h'h'^T * P_i * (J - P_i) - DIAG(h'h'^T * P_i * (J - P_i))\|)$$
$$\leq E(\|2h'h'^T * P_i * (J - P_i)\|)$$
$$\leq E(\|2h'h'^T * P_i\|)$$
$$\leq 2E(\lambda_i)\|h'h'^T * h_1 h_1^T\|$$
$$= 2E(\lambda_i)$$

Notice that the only non-zero eigenvector of $E(\lambda_i^2)(h_1^T h')^2 h_1 h_1^T$ is h_1 and the corresponding eigenvalue is $E(\lambda_i^2)(h_1^T h')^2$. We apply Davis-Kahan theorem [24] to the eigenvector corresponding to the largest eigenvalue of matrices $E(\langle A_i, h'h'^T \rangle A_i)$ and $E(\lambda_i^2)(h_1^T h')^2 h_1 h_1^T$,

$$\|h' - h_1\| \leq \frac{2E(\lambda_i)}{E(\lambda_i^2)(h_1^T h')^2}$$

CCI Linear Model Summary

```
> model<-lm( cci ~Lambda+1)
> summary(model)

Call:
lm(formula = cci ~ Lambda + 1)

Residuals:
Min        1Q      Median     3Q       Max
-26.3432  -6.144  -0.7578    7.1032   16.9004


Coefficients:  Estimate    Pr(>|t|)
(Intercept)    1.275e+02   0.000275  ***
Lambda1        2.421e-04   0.997981
Lambda2       -2.326e-01   0.070110  .
Lambda3       -3.716e-02   0.822592
Lambda4        8.049e-02   0.687628
Lambda5       -2.925e-01   0.421858
Lambda6       -4.285e-01   0.009088  **
Lambda7       -1.745e-01   0.590533
Lambda8       -3.465e-01   0.240093
Lambda9       -8.970e-01   0.007999  **
Lambda10      -8.955e-01   0.052839  .
---

Signif. codes:  0    ***     0.001     **
0.01    *    0.05      .     0.1          1

Residual standard error: 9.437
on 102 degrees of freedom
Multiple R-squared:  0.2325,
Adjusted R-squared:  0.1572
F-statistic:  3.09 on 10 and 102 DF,
p-value: 0.001795
```

REFERENCES

[1] E. Otte and R. Rousseau, "Social network analysis: a powerful strategy, also for the information sciences," *Journal of information Science*, vol. 28, no. 6, pp. 441–453, 2002.

[2] R. Govindan and H. Tangmunarunkit, "Heuristics for internet map discovery," in *INFOCOM 2000. Nineteenth Annual Joint Conference of the IEEE Computer and Communications Societies. Proceedings. IEEE*, vol. 3. IEEE, 2000, pp. 1371–1380.

[3] E. T. Bullmore and D. S. Bassett, "Brain graphs: graphical models of the human brain connectome," *Annual review of clinical psychology*, vol. 7, pp. 113–140, 2011.

[4] M. D. Ward, K. Stovel, and A. Sacks, "Network analysis and political science," *Annual Review of Political Science*, vol. 14, pp. 245–264, 2011.

[5] I. Jolliffe, *Principal component analysis*. Wiley Online Library, 2002.

[6] G. Li, M. Semerci, B. Yener, and M. J. Zaki, "Graph classification via topological and label attributes," in *Proceedings of the 9th international workshop on mining and learning with graphs (MLG), San Diego, USA*, vol. 2, 2011.

[7] J. Huan, W. Wang, and J. Prins, "Efficient mining of frequent subgraphs in the presence of isomorphism," in *Data Mining, 2003. ICDM 2003. Third IEEE International Conference on*. IEEE, 2003, pp. 549–552.

[8] Z. Zhao and H. Liu, "Spectral feature selection for supervised and unsupervised learning," in *Proceedings of the 24th international conference on Machine learning*. ACM, 2007, pp. 1151–1157.

[9] D. L. Sussman, M. Tang, D. E. Fishkind, and C. E. Priebe, "A consistent adjacency spectral embedding for stochastic blockmodel graphs," *Journal of the American Statistical Association*, vol. 107, no. 499, pp. 1119–1128, 2012.

[10] B. Karrer and M. E. Newman, "Stochastic blockmodels and community structure in networks," *Physical Review E*, vol. 83, no. 1, p. 016107, 2011.

[11] S. J. Young and E. R. Scheinerman, "Random dot product graph models for social networks," in *Algorithms and models for the web-graph*. Springer, 2007, pp. 138–149.

[12] J. D. Carroll and J.-J. Chang, "Analysis of individual differences in multidimensional scaling via an n-way generalization of eckart-young decomposition," *Psychometrika*, vol. 35, no. 3, pp. 283–319, 1970.

[13] W. Tang, Z. Lu, and I. S. Dhillon, "Clustering with multiple graphs," in *Data Mining, 2009. ICDM'09. Ninth IEEE International Conference on*. IEEE, 2009, pp. 1016–1021.

[14] T. G. Kolda, "Numerical optimization for symmetric tensor decomposition," *Mathematical Programming*, vol. 151, no. 1, pp. 225–248, 2015.

[15] J. Nocedal and S. Wright, *Numerical optimization*. Springer Science & Business Media, 2006.

[16] A. Athreya, C. Priebe, M. Tang, V. Lyzinski, D. Marchette, and D. Sussman, "A limit theorem for scaled eigenvectors of random dot product graphs," *Sankhya A*, pp. 1–18, 2013. 2016

[17] J. Neyman and E. L. Scott, "Consistent estimates based on partially consistent observations," *Econometrica: Journal of the Econometric Society*, pp. 1–32, 1948.

[18] M. Belkin and P. Niyogi, "Laplacian eigenmaps for dimensionality reduction and data representation," *Neural computation*, vol. 15, no. 6, pp. 1373–1396, 2003.

[19] R. Arden, R. S. Chavez, R. Grazioplene, and R. E. Jung, "Neuroimaging creativity: a psychometric view," *Behavioural brain research*, vol. 214, no. 2, pp. 143–156, 2010.

[20] R. S. Desikan, F. Ségonne, B. Fischl, B. T. Quinn, B. C. Dickerson, D. Blacker, R. L. Buckner, A. M. Dale, R. P. Maguire, B. T. Hyman *et al.*, "An automated labeling system for subdividing the human cerebral cortex on mri scans into gyral based regions of interest," *Neuroimage*, vol. 31, no. 3, pp. 968–980, 2006.

[21] T. M. Amabile, "The social psychology of creativity: A componential conceptualization." *Journal of personality and social psychology*, vol. 45, no. 2, p. 357, 1983.

[22] R. I. Jennrich, "Asymptotic properties of non-linear least squares estimators," *The Annals of Mathematical Statistics*, vol. 40, no. 2, pp. 633–643, 1969.

[23] P. J. Bickel and K. A. Doksum, *Mathematical Statistics: Basic Ideas and Selected Topics, volume I*. CRC Press, 2015, vol. 117.

[24] C. Davis and W. M. Kahan, "The rotation of eigenvectors by a perturbation. iii," *SIAM Journal on Numerical Analysis*, vol. 7, no. 1, pp. 1–46, 1970.